

ABHINAV ANAND

New Delhi, India | +91 9958044538 | abhinava676@gmail.com | linkedin.com/in/abhinava-anand | github.com/kaijin07

TECHNICAL SKILLS

AI & Agents: LangGraph.js, multi-agent orchestration, function calling, prompt engineering, RAG, embeddings, chunking, vector search, reranking, MCP servers, Whisper STT, Groq / LLaMA 3.3, Vapi

Languages: JavaScript (ES6+), TypeScript, Python, SQL, C++, Java

Backend & APIs: Node.js, Express, FastAPI, Flask, REST, WebSockets, Socket.IO, WebRTC, signed webhooks, OAuth 2.0, JWT (HS256), Redis rate limiting, CSP

Frontend: React 19, Next.js 15, Tailwind CSS, shadcn/ui

Data & Infrastructure: PostgreSQL, Supabase (pgvector), MySQL, MongoDB, Redis, schema design, AWS EC2, Docker, n8n, Vercel, Git/GitHub

EXPERIENCE

PrismaVoices Global LLP

May 2026 – Present

Software Engineering Intern | Next.js 15, TypeScript, FastAPI, PostgreSQL, Redis, Vapi

Remote

- Shipped the homepage voice-agent demo in React/TypeScript against a WebRTC voice SDK, and closed the acquisition loop by converting completed calls into qualified pre-signup leads through a signed end-of-call webhook written idempotently to PostgreSQL on call ID.
- Kept long-lived voice credentials out of the browser by minting 10-minute scoped HS256 JWTs server-side in FastAPI, rate-limited to two sessions per IP per day via Redis.
- Fixed dashboard org switching, where the sidebar and API routes could resolve different organizations on one page, by persisting the selection in a cookie authorized against `org_members` on every read (forged IDs fail closed) and replacing the unordered LIMIT 1 membership lookup with a deterministic fallback.

PROJECTS

Recall, Multi-Agent Research Assistant | *Node.js, Express, LangGraph.js, Groq (LLaMA 3.3), Supabase pgvector, Next.js 15, MCP, n8n, AWS* [GitHub](#)

- Orchestrated a LangGraph.js planner that decomposes a question into sub-questions, fans them out to parallel searcher agents holding isolated contexts and distinct tool sets, and synthesizes one cited answer, logging every tool call to PostgreSQL for replay.
- Grounded answers in a personal document corpus with a RAG pipeline on Supabase pgvector: fixed-window chunking, local Transformers.js embeddings, top-*k* retrieval and cross-encoder reranking, with retrieved chunks surfaced in the UI for provenance.
- Served one Express agent core to three clients: a Next.js dashboard, a Telegram bot handling voice notes via Groq whisper-large-v3, and an MCP server exposing the corpus to Claude Desktop; n8n ingests on a schedule, admitting only documents whose embeddings match user-defined topics. Deployed on AWS EC2.

Hermes AI, AI Customer Support Platform | *React 19, Node.js, Groq (LLaMA 3.1), MongoDB, Socket.IO, Google OAuth, JWT* [GitHub](#) | [Demo](#)

- Resolved repeat support questions from ticket history and PDF knowledge bases with a RAG pipeline: document chunking, embeddings, and top-*k* semantic retrieval passed as context to LLaMA 3.1 via the Groq API.
- Handed live chats from AI to human agents with no context loss using bidirectional Socket.IO events that transfer the full conversation trail.
- Shipped a single-script-tag embeddable widget secured with Google OAuth 2.0, JWT sessions, and Helmet, deployed on Render.

JobGenie, AI-Powered Job Recommendation System | *Python, Flask, NLP, scikit-learn, Groq API, REST APIs* [IEEE Paper](#)

- Achieved a 0.89 F1 score on multi-class job-role classification by building an NLP pipeline over 10,000+ job descriptions with resume parsing, deduplication, and stop-word filtering.
- Delivered real-time ranked recommendations through a Flask REST API backed by Groq LLM inference; won **Best Paper** at EmergIN 2024 and published in **IEEE** (April 2025).

EDUCATION

Chandigarh University

2021 – 2025

Bachelor of Engineering, Computer Science